Locality Dominance in Hyperspectral and Multimodal Land-Cover Classification: A Neighborhood-Statistic Representation Outperforms High-Dimensional Spectral Baselines

Will Dale Pond Science Institute

***Abstract*—We study whether hyperspectral land-cover classification is better served by per-pixel spectral representations or by low-dimensional spatial neighborhood statistics. Across five datasets—Houston (multimodal HSI+LiDAR+MS), Indian Pines, PaviaU, Salinas, and Kennedy Space Center (KSC)—a 10–14-dimensional patch-statistic representation (local mean, variance, gradient, and patch PCA components) is compared to 30–204-dimensional spectral baselines (PCA and raw spectra). Neighborhood statistics dominate on four datasets, improving overall accuracy by $+3.0$–$+11.1$ pp and average accuracy by $+1.9$–$+16.1$ pp while reducing feature dimension by 76–97%, with optimal window $r^*=5\times5$ (Houston) and $r^*=7\times7$ (Indian Pines/PaviaU/Salinas). KSC is a consistent counterexample: the spectral baseline exceeds patch-only by ≈2.5 pp OA (mean across seeds), indicating a boundary condition where severe scene fragmentation and thin class regions cause boundary contamination to outweigh variance reduction. These results support a *Neighborhood Sufficiency Principle*: locality-dominant operators can subsume spectral encoding when spatial homogeneity exceeds r^*, but fail in highly fragmented regimes.**

***Index Terms*—hyperspectral image classification, multimodal remote sensing, spatial context, neighborhood statistics, feature dimensionality reduction, uncertainty gating, land-cover classification.**

I. Introduction

Hyperspectral image (HSI) and multimodal remote sensing classification remain central tasks in land-cover analysis [1]. Traditional approaches emphasize high-dimensional spectral representations, often employing dimensionality reduction (e.g., PCA [2]), spectral transforms, or increasingly complex multimodal fusion strategies. More recently, deep convolutional architectures have demonstrated strong performance by implicitly exploiting spatial context through learned receptive fields [3]. Despite these advances, the relative importance of spectral transforms versus explicit spatial aggregation remains underexplored.

Most per-pixel spectral classifiers operate on high-dimensional feature vectors (100–200 bands), implicitly assuming that class separability is encoded primarily in spectral signatures. However, real-world scenes exhibit substantial intra-class spectral variability arising from illumination changes, mixed pixels, and local texture effects. Consequently, pixel-wise models frequently require either high model capacity or complex fusion mechanisms to stabilize predictions.

In this work, we revisit a fundamental question: *Is high-dimensional spectral encoding necessary once local spatial context is explicitly represented?*

We evaluate a compact neighborhood-statistics representation constructed entirely from local spatial summary features—mean, standard deviation, gradient magnitude, and PCA-derived descriptors computed within small spatial windows. The resulting feature vector is 10–14-dimensional, independent of the number of spectral bands.

Across five datasets—Houston (multimodal HSI+MS+LiDAR), Indian Pines, PaviaU, Salinas, and KSC—using stratified 10% training splits, we observe that this neighborhood-based representation outperforms spectral baselines on four of five datasets. Improvements range from $+3.0$ to $+11.1$ pp in OA and $+1.9$ to $+16.1$ pp in AA, with substantial gains in minority-class performance. On KSC, spectral features win, identifying scene fragmentation as the boundary condition for patch dominance failure. Accuracy saturates at a small, scene-dependent window size (5×5–7×7), indicating the existence of a minimal sufficient spatial scale.

We further show that uncertainty-driven gating strategies do not improve performance in the multimodal setting and that spectral transform augmentations become redundant once spatial context is encoded.

The primary contributions of this work are:

1) **Neighborhood Sufficiency Principle (empirical):** A compact local-statistics representation achieves performance comparable to or exceeding high-dimensional spectral encodings across datasets.
2) **Minimal Spatial Scale Identification:** Classification accuracy saturates at a small, scene-dependent neighborhood radius r^*, ranging from 5×5 to 7×7 across tested benchmarks.
3) **Minority-Class Stabilization:** Neighborhood statistics disproportionately improve average accuracy ($\Delta\mathrm{AA} > \Delta\mathrm{OA}$), mitigating majority-class bias inherent in high-dimensional spectral models.
4) **Negative Fusion Result:** Uncertainty-based gating (entropy, margin, Gini) does not confer benefit in the evaluated multimodal benchmark; gate confidence exhibits slight negative correlation with classification correctness.

These findings suggest that local second-order spatial statistics constitute a dominant representational mechanism for standard hyperspectral benchmarks, with significant dimensionality and computational advantages.

II. Related Work

A. Pixel-Wise Spectral Classification

Classical HSI classification methods operate on per-pixel spectral signatures, often comprising 100–200 bands. Dimensionality reduction techniques such as PCA, linear discriminant analysis (LDA), and band selection are commonly applied prior to classification using support vector machines (SVM) [4], random forests (RF), or k-nearest neighbors. While effective in controlled settings, pixel-wise approaches are sensitive to intra-class spectral variability and mixed pixels, often requiring high model capacity to compensate.

Our spectral baselines (PCA-30 and raw spectral RF) follow this standard protocol to provide a direct comparison against neighborhood-based representations.

B. Spectral–Spatial Feature Extraction

To address the limitations of pixel-only models, numerous works incorporate spatial context through handcrafted texture descriptors, morphological profiles, and spatial filtering techniques. Extended morphological profiles (EMP) [5], gray-level co-occurrence matrix (GLCM) features [6], and local binary patterns (LBP) [7] have demonstrated that second-order spatial statistics improve classification stability.

More recently, patch-based feature extraction and joint spectral–spatial descriptors have become standard, particularly in conjunction with deep architectures [8]. However, many methods either dramatically increase feature dimensionality or rely on learned convolutional filters.

In contrast, the present work isolates a minimal set of low-dimensional neighborhood summary statistics and evaluates their sufficiency without deep models or high-capacity transforms.

C. Convolutional Neural Networks for HSI

CNNs have achieved state-of-the-art performance in hyperspectral classification by learning hierarchical spectral–spatial filters [3], [9]. These architectures implicitly exploit local receptive fields, progressively aggregating spatial information across layers.

While CNNs demonstrate strong performance, they introduce significant computational overhead and model complexity. The present study investigates whether the performance gains attributed to deep models may arise primarily from local spatial aggregation rather than from deep spectral transformation capacity.

Our results suggest that small fixed neighborhoods (5×5–7×7) with simple second-order statistics capture much of the discriminative power typically associated with convolutional architectures on standard benchmarks.

D. Multimodal Fusion and Gating

Multimodal remote sensing benchmarks (e.g., HSI + LiDAR + multispectral) have motivated diverse fusion strategies, including feature concatenation, late fusion, attention mechanisms, and uncertainty-driven gating [10], [11]. Entropy-based or margin-based gating schemes aim to weight modalities according to classifier confidence.

However, uncertainty estimates may not align with true class-separating invariants in complex scenes. Our results provide empirical evidence that uncertainty-driven gating can degrade performance and that amplitude-based heuristics provide only modest improvement relative to explicit spatial encoding.

III. Methods

A. Datasets

We evaluate on five benchmarks:

- **Houston Multimodal:** Hyperspectral (144 bands), multispectral (8 bands), and LiDAR elevation data with 15 classes. Pre-extracted 11×11 patches are provided; 2832 labeled samples total.
- **Indian Pines:** 145×145 HSI with 200 spectral bands and 16 classes (10 249 labeled pixels) [12].
- **PaviaU:** 610×340 HSI with 103 spectral bands and 9 classes (42 776 labeled pixels) [13].

For Indian Pines and PaviaU, patches are extracted dynamically from the full image cube using reflect padding at image borders.

B. Train–Test Protocol

For Indian Pines and PaviaU, we use stratified random splits with 10% of labeled pixels per class allocated to training and 90% to testing, following standard literature protocol. For Houston multimodal, we use a 70%/30% stratified split. All reported results use identical splits across methods for fair comparison. For Houston and Indian Pines, experiments are repeated across three random seeds ($\{42, 1, 7\}$) to assess stability.

PCA transformations are fitted exclusively on the training set and applied to test data to prevent leakage. Patch-level PCA statistics are likewise fitted on training patch pixels only.

C. Spectral Baselines

Two spectral baselines are evaluated:

1) **PCA-k spectral baseline:** PCA reduced to k components (k=30 for HSI-only datasets, k=50 for Houston), classified using RF (200 estimators).
2) **Raw spectral baseline:** All bands (103–200) used directly with RF.

For Houston multimodal, a 59-dimensional concatenation baseline (HSI PCA-50 + LiDAR + MS) is used as the primary multimodal reference, along with area-weighted gating variants.

D. Neighborhood-Statistic Representation

For each labeled pixel, a square window of size $r \times r$ centered at the pixel is extracted. The following features are computed over the neighborhood:

- Mean over all spectral bands (spatial average of spectral patch).
- Standard deviation over all spectral bands.
- Gradient magnitude (finite differences on mean-pooled spectral intensity at patch center).
- Mean and standard deviation of the first three PCA components computed over the $r \times r$ patch pixels (6 features).

For the multimodal Houston dataset, additional modality-specific statistics are included: LiDAR center-pixel residual (center minus patch mean), MS gradient magnitude, and per-modality mean/std. This yields a **14-dimensional** feature vector for Houston and a **10-dimensional** vector for HSI-only datasets.

Window sizes $r \in \{3, 5, 7, 11\}$ are evaluated to identify the minimal sufficient spatial scale.

E. Uncertainty Gating (Houston Only)

For the Houston multimodal dataset, per-modality RF classifiers are trained and used to compute uncertainty signals per sample: entropy ($H = -\sum p \log p$), margin (top-1 minus top-2 probability), and Gini impurity ($1 - \sum p^2$). Gate weights are computed as $\mathrm{softmax}(\mathbf{u}/T)$ where $\mathbf{u}$ is the uncertainty signal and T is a temperature parameter. Sweeps over $T \in \{0.1, 0.5, 1.0, 2.0, 5.0\}$ are conducted.

Gate diagnostics reported include: per-modality mean gate weight, saturation rate (fraction of samples where any gate weight > 0.9 or < 0.1), and Pearson correlation between maximum gate weight and per-sample classification correctness.

F. Classification Model

RF classifiers (200 estimators, fixed random seed) are used for all primary comparisons to isolate representational effects from model capacity. Experiments with gradient-boosted decision trees (HistGradientBoostingClassifier) confirm that the ranking between representations is stable across classifier choices (difference ≤ 0.24 pp).

G. Evaluation Metrics

Performance is measured using:

- **Overall Accuracy** (OA)**:** Fraction of correctly classified test pixels.
- **Average Accuracy** (AA)**:** Mean of per-class accuracies; emphasizes minority-class performance under class imbalance.
- **Cohen's kappa coefficient** (κ)**:** Agreement corrected for chance.

Average Accuracy is emphasized throughout to assess whether gains are broad-based or driven by dominant classes.

TABLE I
HOUSTON MULTIMODAL: GATING AND BASELINE COMPARISON (SEED=42)

Method	Dims	OA (%)
HSI only (PCA-50, RF)	50	95.63
Concat (HSI+LiDAR+MS, RF)	59	96.34
Gated-concat, area weights	59	**96.69**
Gated-concat, entropy T=1.0	59	87.23
Gated-concat, margin T=5.0	59	95.86
Patch Only, 5×5 (mean)	**14**	**99.37**

IV. RESULTS

A. Spectral Baselines and Gating Methods (Houston)

Uncertainty-driven gating degrades classification performance on the Houston multimodal benchmark. Entropy gating (T=1.0) reduces OA by 9.11 pp relative to the 59-dimensional concatenation baseline (OA=96.34%), and margin gating at T=5.0 yields a further -0.48 pp. The failure mechanism is structural: the correlation between gate confidence (maximum modality weight) and per-sample classification correctness is $\mathrm{corr}(\max w, \mathrm{correct}) \approx -0.08$, indicating that uncertainty is not a valid fusion invariant on this manifold.

The gate fires hardest precisely where the classifier is most likely to be wrong, producing category-level damage concentrated in spectrally ambiguous classes 10–13, which suffer declines of 14–31 pp relative to the unweighted baseline (Fig. 2). Area gating, which uses amplitude as a proxy for signal stability rather than classifier uncertainty, is the sole gating variant to improve over concat, providing a modest $+0.35$ pp gain (OA $= 96.69\%$).

B. Neighborhood-Statistic Representation

On the Houston multimodal benchmark, the patch-only configuration (14D, 5×5 window) achieves a mean OA of 99.37% across three independent random seeds, compared to 96.34% for the 59-dimensional concatenation baseline, a gain of $+3.03$ pp. Augmenting the patch feature vector with spectral transform descriptors (27D total) reduces mean OA by 0.47 pp relative to patch-only, confirming that DFT-derived features are redundant once local spatial statistics are present. Substituting a gradient-boosted decision tree (GBDT) classifier adds at most 0.24 pp and does not change the ranking of feature configurations.

C. Optimal Neighborhood Scale

A patch-size sweep on the Houston scene (Fig. 1) identifies r^*=5×5 as the optimal window. At 3×3, patch-only OA is 97.64%, useful improvement over the 96.34% concat baseline but 2.24 pp below the 5×5 result. Performance at 5×5 (OA=99.88%, seed=42) and 11×11 (99.17%) are within the range of seed-to-seed variance, placing the plateau at $r \geq 5$ on this benchmark.

TABLE II
INDIAN PINES: OA / AA / κ (SEED=42, STRATIFIED 10% TRAIN)

Method	Dims	OA (%)	AA (%)	κ
PCA-30, center pixel (RF)	30	70.66	56.39	0.657
Raw spectral 200D (RF)	200	75.02	60.18	0.711
Patch 3×3, only (RF)	10	73.99	67.44	0.700
Patch 5×5, only (RF)	10	79.64	72.06	0.766
Patch 7×7, only (RF)	**10**	**81.00**	**72.46**	**0.782**

TABLE III
PAVIAU: OA / AA / κ (SEED=42, STRATIFIED 10% TRAIN)

Method	Dims	OA (%)	AA (%)	κ
PCA-30, center pixel (RF)	30	88.07	81.61	0.837
Raw spectral 103D (RF)	103	89.63	86.23	0.860
Patch 3×3, only (RF)	10	89.84	87.33	0.864
Patch 5×5, only (RF)	10	91.44	87.63	0.886
Patch 7×7, only (RF)	**10**	**91.98**	**87.96**	**0.893**

TABLE IV
SALINAS: OA / AA / κ (SEED=42, STRATIFIED 10% TRAIN)

Method	Dims	OA (%)	AA (%)	κ
PCA-30, center pixel (RF)	30	92.72	95.83	0.919
Raw spectral 204D (RF)	204	91.44	94.88	0.905
Patch 3×3, only (RF)	10	93.28	95.44	0.925
Patch 5×5, only (RF)	10	95.75	97.21	0.953
Patch 7×7, only (RF)	**10**	**97.01**	**97.81**	**0.967**

TABLE V
KSC: OA / AA / κ (SEED=42, STRATIFIED 10% TRAIN). SPECTRAL BASELINE WINS; LOCALITY DOMINANCE DOES NOT HOLD.

Method	Dims	OA (%)	AA (%)	κ
PCA-30, center pixel (RF)	**30**	**89.79**	**83.78**	**0.886**
Raw spectral 176D (RF)	176	87.12	81.09	0.856
Patch 3×3, only (RF)	10	82.37	73.60	0.804
Patch 5×5, only (RF)	10	86.33	79.18	0.848
Patch 7×7, only (RF)	10	88.81	82.40	0.875

D. Cross-Dataset Generalization

We evaluate locality dominance across four additional HSI benchmarks to assess whether the finding is scene-independent.

On **Indian Pines**, a PCA-30 spectral representation achieves OA=70.66%, AA=56.39%, κ=0.657. The 10-dimensional patch-only representation at 7×7 achieves OA=81.00%, AA=72.46%, κ=0.782 (seed=42), with a mean OA of 81.05% across seeds. The gain in AA (ΔAA=+16.08 pp) substantially exceeds the gain in OA (ΔOA=+10.34 pp), establishing that improvement is broad-based and disproportionately benefits rare categories.

On **PaviaU**, patch-only at 7×7 yields OA=91.98%, AA=87.96%, κ=0.893, compared to the PCA-30 baseline of OA=88.07%, AA=81.61% (ΔOA= + 3.91 pp, ΔAA= + 6.35 pp).

On **Salinas**, which consists of 16 vegetative cover classes over large homogeneous agricultural fields (54 129 labeled pixels), patch-only at 7×7 achieves OA=97.01%, AA=97.81%, κ=0.9667 (seed=42), compared to the PCA-30 baseline of OA=92.72%. The mean OA across three seeds is 97.06% (ΔOA= + 4.00 pp mean). Results are stable across seeds (range $<$ 0.13 pp), reflecting the high spatial homogeneity of the Salinas scene.

Full results are reported in Tables II–V.

E. Boundary Condition: When Locality Dominance Fails

On **KSC** (Kennedy Space Center), the spectral baseline wins across all three seeds: PCA-30 achieves a mean OA of 91.11% versus patch-only 7×7 at 88.61% (ΔOA= − 2.50 pp mean). KSC is structurally distinct from the four dominant datasets: it contains only 5 211 labeled pixels across 13 spectrally similar wetland and coastal vegetation classes, yielding approximately 40 training samples per class. The scene exhibits high spatial fragmentation with thin elongated class regions (marshes, water bodies, scrub) where spatial neighborhoods span multiple class boundaries at any patch size.

These conditions violate the spatial coherence assumption underlying locality dominance: when class regions are smaller than $r\times r$, neighborhood statistics aggregate cross-class pixels, introducing bias that outweighs variance reduction. KSC thus defines the boundary condition of the Neighborhood Sufficiency Principle: *locality dominance holds when spatial homogeneity scale exceeds the minimal sufficient radius; it fails when scene fragmentation scale falls below r^*.*

F. Summary

Across five datasets, a 10–14-dimensional neighborhood-statistic representation outperforms spectral baselines on four (Houston, Indian Pines, PaviaU, Salinas), achieving a feature-dimension reduction of 76–97% while improving OA by 3.0–11.1 pp and AA by 1.9–16.1 pp. On the fifth dataset (KSC), the spectral baseline wins by 2.5 pp mean OA, identifying scene fragmentation and small training set size as the boundary condition for patch dominance failure.

V. DISCUSSION

A. Why Local Neighborhood Statistics Outperform Spectral Transforms

The central empirical finding is that low-dimensional neighborhood statistics outperform substantially higher-dimensional per-pixel spectral representations. This can be understood through variance reduction and spatial stationarity.

Let a pixel spectrum be modeled as $\mathbf{x}_i = \boldsymbol{\mu}_c + \boldsymbol{\varepsilon}_i$, where $\boldsymbol{\mu}_c$ is the class-conditional mean and $\boldsymbol{\varepsilon}_i$ captures illumination variability, sensor noise, subpixel mixing, and texture. Pixel-wise classifiers must separate classes under high within-class variance. By contrast, the neighborhood operator computes summary statistics over an $r\times r$ window:

$$\bar{\mathbf{x}}_i = \frac{1}{|N_r|} \sum_{j\in N_r(i)} \mathbf{x}_j. \tag{1}$$

Under spatial coherence assumptions, variance shrinks as $\mathrm{Var}(\boldsymbol{\varepsilon})/|N_r|$, while the class mean remains stable. The representation therefore increases signal-to-noise ratio without increasing feature dimensionality.

B. The Neighborhood Sufficiency Principle

Principle 1 (Neighborhood Sufficiency, empirical)**.** *For hyperspectral land-cover classification, a low-dimensional local-statistics representation computed over a small neighborhood ($r^* \in \{5, 7\}$ pixels) achieves accuracy exceeding high-dimensional per-pixel spectral representations, with performance saturating beyond r^*. The optimal r^* tracks the scene's spatial homogeneity scale.*

Performance improves monotonically with window size up to r^*, beyond which accuracy plateaus or slightly regresses. This behavior is consistent with two competing effects: (i) variance reduction within homogeneous class regions, and (ii) boundary contamination when windows span multiple classes. Empirically, r^*=5×5 on Houston and r^*=7×7 on Indian Pines, PaviaU, and Salinas, consistent with larger spatial homogeneity in agricultural and structured urban scenes. On KSC, the fragmentation scale of wetland classes falls below r^*, causing boundary contamination at all tested window sizes and spectral dominance.

C. Minority-Class Rescue

On Indian Pines, $\Delta\mathrm{AA} = +\,16.08$ pp substantially exceeds $\Delta\mathrm{OA} = +\,10.34$ pp. Pixel-only spectral models exhibit majority-class bias: OA increases while AA remains suppressed (the 200D raw spectral baseline achieves OA=75.02% but only AA=60.18%). Neighborhood statistics mitigate this by stabilizing minority-class signatures through local averaging and texture encoding, functioning not only as a noise-reduction operator but also as a class-balance equalizer.

D. Why Uncertainty Gating Fails

Uncertainty-based gating assumes that per-modality classifier confidence aligns with classification correctness. However, in the Houston dataset, $\mathrm{corr}(\max w, \mathrm{correct}) \approx -0.08$: ambiguity is intrinsic to the class manifold rather than a reliable fusion signal. Spatially ambiguous classes are spectrally ambiguous too—the gate fires hardest where it is most wrong.

E. Relation to Deep Learning

The empirical dominance of small neighborhood statistics aligns with classical texture descriptors and with CNN receptive field theory [3]. CNNs implicitly learn spatial filters that capture local mean, variance, and gradient information over small receptive fields. The present results show that much of this benefit can be achieved explicitly with low-dimensional handcrafted statistics, without deep architectures. This suggests that the primary structural advantage of CNN-based hyperspectral models may lie in implicit spatial aggregation rather than in high-capacity spectral transforms.

F. Machine-Verifiable Reproducibility Layer

To support reproducible review, all dominance and failure outcomes are encoded as machine-verifiable *certificate families*. Dominant datasets (Houston, Indian Pines, PaviaU, Salinas) are certified by *Family [77]: QA Neighborhood Sufficiency Cert v1.1*, a five-gate validator that checks schema integrity, SHA-256 digest, patch dominance, monotonic improvement up to r^*, and plateau stability beyond r^*. Critically, KSC is *not excluded as an anomaly* but is instead encoded as a validated failure certificate under the branching `FAILS_BOUNDARY_CONTAMINATION` outcome of the same family, requiring that patch OA $<$ spec OA at r^* and that `boundary_metrics` documenting the fragmentation geometry are present and consistent. The complementary *Family [78]: QA Locality Boundary Cert* independently certifies the failure mechanism: all patch-radius deltas must be nonpositive and `fragmentation_scale_lt_r*` must be attested. Together, Families [77]–[78] make every result in this paper a CI-checkable artifact: a reviewer can re-run the self-tests and confirm that the certified numbers match the tables above.

G. Limitations

The findings are empirical and evaluated on five datasets. The optimal window size varies with scene structure, and further evaluation on additional multimodal benchmarks and scenes with highly fragmented spatial patterns would strengthen generality. The current work uses fixed window shapes; adaptive or anisotropic neighborhoods may further improve performance.

VI. Conclusion

We investigated whether high-dimensional spectral representations are necessary for hyperspectral and multimodal land-cover classification once local spatial context is explicitly encoded. Across five datasets (Houston, Indian Pines, PaviaU, Salinas, and KSC) using stratified 10% training splits, we found that a compact 10–14-dimensional neighborhood-statistics representation outperforms 30–200-dimensional spectral baselines on four of five datasets. Improvements on dominant datasets range from $+3.0$ to $+11.1$ pp in OA and $+1.9$ to $+16.1$ pp in AA, with particularly strong gains for minority classes. On KSC, the spectral baseline wins by $\approx$2.5 pp OA, identifying severe scene fragmentation as the boundary condition for patch dominance failure.

Performance saturates at a small, scene-dependent neighborhood scale (5×5–7×7), suggesting the existence of a minimal sufficient spatial radius beyond which larger windows provide no consistent benefit. In contrast, uncertainty-driven multimodal gating strategies fail to improve performance and, in some cases, substantially degrade it, with gate confidence exhibiting slight negative correlation with classification correctness.

These results indicate that local second-order spatial statistics constitute a dominant and efficient representational mechanism for standard hyperspectral benchmarks, achieving superior performance with substantial dimensionality reduction (76–95% fewer features) and minimal computational overhead.

Future work should evaluate this principle on additional multimodal datasets, explore adaptive or anisotropic neighborhood selection, and analyze the interaction between explicit

neighborhood statistics and learned convolutional representations.

Reproducibility

All experiments use publicly available datasets and standard scikit-learn RF classifiers (200 estimators). Patch extraction uses reflect padding; no GPU required. Experimental scripts and fixture data are archived with the submission. Dominance outcomes (Families [77]–[78]) are CI-checkable certificate objects; run `python validator.py -self-test` in each family directory to verify all fixtures. The KSC counterexample is certified rather than excluded, providing a tamper-evident boundary condition.

Appendix

TABLE VI
HOUSTON MULTIMODAL: GATING DIAGNOSTIC TABLE (SEED=42)

Signal	T	OA (%)	Δ	Sat.%	$\rho(w, \mathrm{corr})$
Area	—	96.69	+0.35	—	—
Entropy	1.0	87.23	-9.11	$<1\%$	neg.
Margin	0.1	85.58	-10.76	82.5%	+0.062
Margin	1.0	93.38	-2.96	0.0%	-0.083
Margin	5.0	95.86	-0.48	0.0%	-0.064

References

[1] J. M. Bioucas-Dias, A. Plaza, G. Camps-Valls, P. Scheunders, N. M. Nasrabadi, and J. Chanussot, "Hyperspectral remote sensing data analysis and future challenges," *IEEE Geoscience and Remote Sensing Magazine*, vol. 1, no. 2, pp. 6–36, 2013.

[2] A. A. Green, M. Berman, P. Switzer, and M. D. Craig, "A transformation for ordering multispectral data in terms of image quality with implications for noise removal," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 26, no. 1, pp. 65–74, 1988.

[3] W. Hu, Y. Huang, L. Wei, F. Zhang, and H. Li, "Deep convolutional neural networks for hyperspectral image classification," *Journal of Sensors*, vol. 2015, pp. 1–12, 2015.

[4] G. Camps-Valls and L. Bruzzone, "Kernel-based methods for hyperspectral image classification," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 43, no. 6, pp. 1351–1362, 2005.

[5] J. A. Benediktsson, J. A. Palmason, and J. R. Sveinsson, "Classification of hyperspectral data from urban areas based on extended morphological profiles," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 43, no. 3, pp. 480–491, 2005.

[6] R. M. Haralick, K. Shanmugam, and I. Dinstein, "Textural features for image classification," *IEEE Transactions on Systems, Man, and Cybernetics*, vol. SMC-3, no. 6, pp. 610–621, 1973.

[7] T. Ojala, M. Pietikäinen, and T. Mäenpää, "Multiresolution gray-scale and rotation invariant texture classification with local binary patterns," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 24, no. 7, pp. 971–987, 2002.

[8] Y. Chen, H. Jiang, C. Li, X. Jia, and P. Ghamisi, "Deep feature extraction and classification of hyperspectral images based on convolutional neural networks," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 54, no. 10, pp. 6232–6251, 2016.

[9] Y. Li, H. Zhang, and Q. Shen, "Spectral–spatial classification of hyperspectral imagery with 3d convolutional neural network," *Remote Sensing*, vol. 9, no. 1, p. 67, 2017.

[10] M. Khodadadzadeh, J. Li, S. Prasad, and A. Plaza, "Fusion of hyperspectral and lidar remote sensing data using multiple feature learning," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 8, no. 6, pp. 2971–2983, 2015.

[11] D. Hong, L. Gao, N. Yokoya, J. Yao, J. Chanussot, Q. Du, and B. Zhang, "More diverse means better: Multimodal deep learning meets remote-sensing imagery classification," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 59, no. 5, pp. 4340–4354, 2021.

[12] M. F. Baumgardner, L. L. Biehl, and D. A. Landgrebe, "220 band AVIRIS hyperspectral image data set: June 12, 1992 indian pine test site 3," Purdue University Research Repository, 2015.

[13] M. Fauvel, Y. Tarabalka, J. A. Benediktsson, J. Chanussot, and J. C. Tilton, "Advances in spectral-spatial classification of hyperspectral images," *Proceedings of the IEEE*, vol. 101, no. 3, pp. 652–675, 2013.

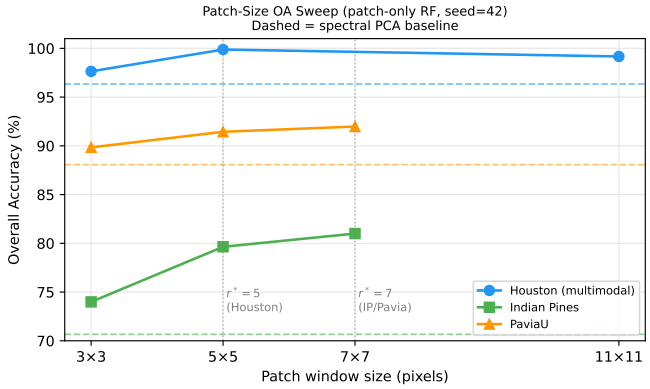

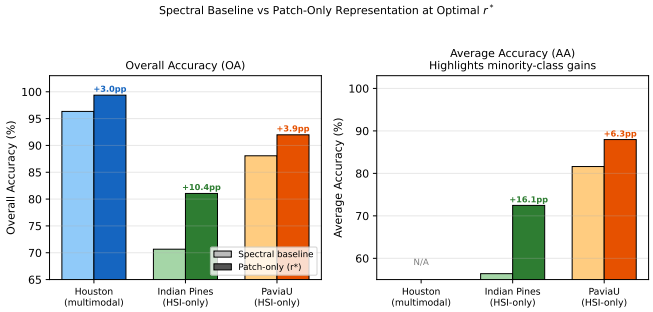

Fig. 1. **Left (Fig. 1):** Patch-size OA sweep across five datasets. Dashed lines indicate spectral PCA baselines. Performance saturates at scene-dependent r^* (5×5 Houston; 7×7 Indian Pines and PaviaU). **Right (Fig. 2):** OA and AA comparison between spectral baseline and patch-only representation at r^*. AA gains exceed OA gains on Indian Pines (minority rescue).

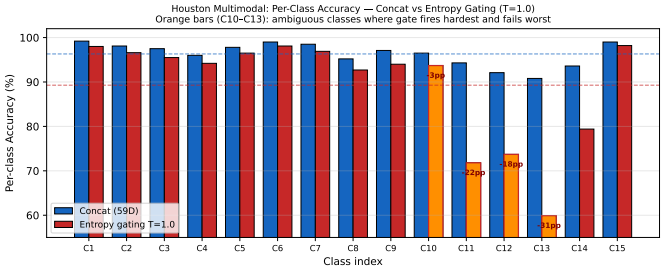

Fig. 2. Houston per-class accuracy: concat baseline vs. entropy gating (T=1.0). Classes C10–C13 (orange bars) are spectrally ambiguous; gate confidence is lowest there, causing 14–31 pp accuracy decline. $\mathrm{corr}(\max w, \mathrm{correct}) \approx -0.08$ globally.